

Ames Housing Challenge

- Antony Paulson Chazhoor






Problem Statement

- Create a regression model based on the Ames Housing Dataset which will predict the price of a house at sale
 - Models are primarily Linear Regression Models, Lasso & Ridge
 - Evaluation Parameter : RMSE
- Why is this analysis important?
 - An analyst can help prospective buyers to understand in cost estimation
 - Overpriced or Underpriced(Mega Deals!!) Houses can be spotted

Exploratory Analysis




Predictors that are highly correlated with Target(Sale Price)

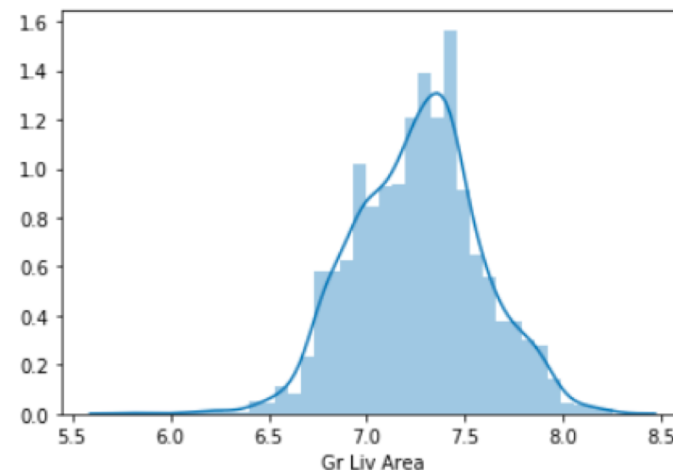







Exploratory Analysis(Contd)

Extremely correlated variables : Garage Area, Garage Cars




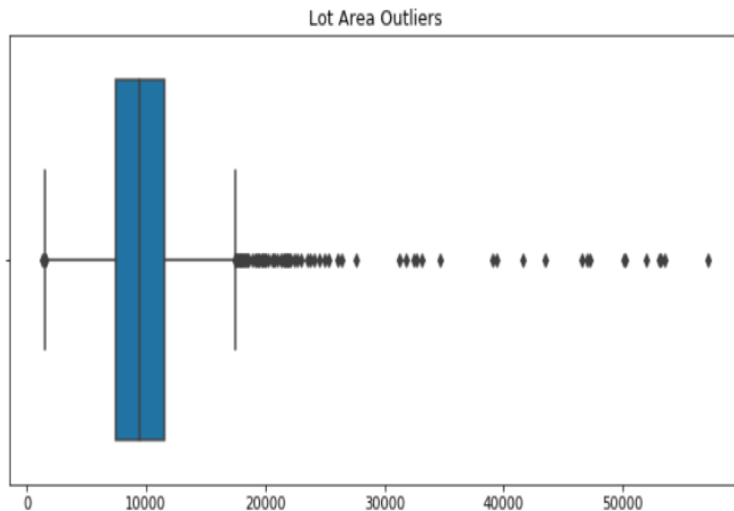
Right Skewed Data



These can be brought to the normal distribution form through log transformations



Extreme Outliers in the Data



Removing these greatly improved RMSE Values



Handling Missing Values

- The Most tricky part in the EDA process
- I used three values to fill my null values:
 - '0' : For numeric data that was missing
 - 'Appropriate String' for categorical predictors
 - 'Forward Fill' : Single missing value.
- Dropped columns with more than 1000 missing values



Preprocessing & Modelling Steps

- One hot encoding of categorical variables was absolutely crucial
- Log transformation of the predicted column SalePrice
- The next step would entail splitting and testing the model
- Scaling the data is necessary for application into Ridge, Lasso models
- First assessment based of the linear regression RMSE
- Cross Validation scores compared across Ridge, Lasso and Linear regression models
- Application of Lasso Model to the Scaled Test data



Creation of Dummy Variables

| | MS Zoning_C (all) | MS Zoning_FV | MS Zoning_I (all) | MS Zoning_RH | MS Zoning_RL | MS Zoning_RM | Street_Pave | Lot Shape_IR2 | Lot Shape_IR3 |
|---|-------------------------|-----------------|-------------------------|-----------------|-----------------|-----------------|-------------|------------------|------------------|
| 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 |
| 2 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 |

- The number of categorical variables: 37
- Number of dummy columns created : 214

Comparison of Cross validation Scores

```
In [1947]: 1 #Checking the cross_val score for Lasso  
          2 cross_val_score(lasso, X_train_sc, y_train, cv = 5).mean()
```

```
Out[1947]: 0.9108505557495308
```

```
In [1948]: 1 ##Checking the cross_val score for LR  
          2 cross_val_score(lr, X_train, y_train, cv =5).mean()
```

```
Out[1948]: 0.9019243802868859
```

```
In [1949]: 1 ##Checking the cross_val score for Ridge  
          2 cross_val_score(ridge, X_train_sc, y_train, cv =5).mean()
```

```
Out[1949]: 0.9022538595296308
```



Results

- The Lasso model from sklearn provided the best accuracy among models
- About 91% of change in the Sales Price could be accounted for by the variables in my model.
- The most correlated columns to sales price were quality & Living area
- The lowest RMSE obtained by the model was approx. 19243.









Future Improvements

- The model can be improved further by implementing inferences from the EDA:
 - Log transformations
 - Dropping highly inter correlated columns
 - Excluding more outliers
 - Eliminating predictors which have very little effect on sale Price



Results of Predictions

| # | Team Name | Kernel | Team Members | Score ? | Entries | Last |
|---|----------------|--------|---|--------------|---------|------|
| 1 | DTrichter | |  | 18630.727... | 46 | 21m |
| 2 | Laura Luo | |  | 18682.682... | 10 | 9h |
| 3 | Nick Minaie | |  | 18860.115... | 52 | 8h |
| 4 | minion_of_boom | |  | 19059.306... | 10 | 1d |
| 5 | Joey Romness | |  | 19075.814... | 55 | 14h |
| 6 | Tony | |  | 19243.121... | 19 | now |

Your Best Entry [↑](#)

Your submission scored 19545.28111, which is not an improvement of your best score. Keep trying!